

## Title: Exploring Machine Learning Concepts with Logistic Regression

**Introduction** In the fascinating world of machine learning, concepts such as data preprocessing, visualization, and model building form the foundation of effective predictive systems. Today, let's unravel these interrelated ideas through a real-world example that classifies rocks and mines using logistic regression. Whether you're a beginner or a seasoned practitioner, understanding these steps can deepen your appreciation for machine learning algorithms.

**The Dataset** We begin with the Sonar dataset, which consists of 208 rows and 61 columns. The dataset includes numerical features that represent sonar signals, with the last column indicating the label: **'R' for Rock** and **'M' for Mine**. This binary classification problem is an excellent use case for logistic regression.

**Step 1: Libraries** To handle data efficiently, we use Python libraries such as:

- **Pandas** for data manipulation.
- **NumPy** for numerical operations.
- **Seaborn** and **Matplotlib** for visualization.
- **Scikit-learn** for machine learning tools.

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from sklearn.linear_model import LogisticRegression
```

**Step 2: Understanding the Dataset** The initial steps involve loading the data, exploring its structure, and performing statistical summaries.

```
df = pd.read_csv('Sonar_data.csv')
print(df.head())
print(df.info())
print(df.describe())
```

Here, we observe the distribution of labels and ensure there are no missing values. Visualizations like bar plots provide insights into class imbalances:

```
sns.barplot(x=df['61'].unique(), y=df['61'].value_counts())
plt.title("Count of Rock and Mines")
plt.show()
```

**Step 3: Data Preprocessing** We separate the features (columns 1-60) from the target label (column 61):

```
x = df.drop(columns='61', axis=1)
y = df['61']
```

The data is then split into training and test sets:

```
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.5, stratify=y, random_state=1)
```

**Step 4: Model Training** Logistic Regression, a simple yet powerful algorithm for binary classification, is used here. The model is trained using the training data:

```
model = LogisticRegression()
model.fit(x_train, y_train)
```

**Step 5: Model Evaluation** Finally, the model is evaluated on test data to check its accuracy:

```
y_pred = model.predict(x_test)
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy}")
```

### Key Concepts at Play

1. **Data Visualization:** Helps in understanding the dataset and identifying patterns.
2. **Train-Test Split:** Ensures the model generalizes well to unseen data.
3. **Logistic Regression:** A foundational algorithm for binary classification problems.
4. **Performance Metrics:** Accuracy is a common metric for evaluating classification models.

**Conclusion** This hands-on example demonstrates the interconnected steps in building a machine learning model. From preprocessing to evaluation, every step contributes to the model's predictive power. Logistic Regression, despite its simplicity, serves as an excellent starting point for understanding machine learning algorithms.